

Future-Proofing Financial AI: A Decentralized Approach to Quantum-Resistant Federated Learning

Pranav Manoj P
Department of Computer Science
PES University
Electronic city, Bengaluru

Chandana Lakya
Department of Computer Science
PES University
Electronic city, Bengaluru

Dr. P Kavitha
Department of Computer Science
PES University
Electronic city, Bengaluru

Abstract—In today’s digital banking environment, protecting sensitive consumer data is paramount as institutions increasingly rely on AI-driven models for credit risk assessment. Traditional encryption methods like RSA and ECC are vulnerable to attacks from emerging quantum computers, which can break these schemes in polynomial time using Shor’s algorithm. To address this, we propose a federated learning framework enhanced with decentralized post-quantum cryptography (PQC-FL) to futureproof banking systems. Our system integrates a hybrid cryptographic scheme that combines Kyber for quantum-safe key exchange, Dilithium and Falcon for quantum-safe digital signatures, and AES for efficient symmetric encryption. These cryptographic primitives were selected in accordance with NIST’s postquantum cryptography standards to provide quantum-resilient confidentiality, integrity, and authentication. Unlike typical federated learning setups that rely on a single central server, our method uses a decentralized architecture with three servers that collaboratively coordinate model training, thereby enhancing fault tolerance and trust by avoiding a single point of failure. In our framework, banking clients train local models on non-IID customer data and share only encrypted model updates. These updates are securely aggregated by decentralized servers using protocols that preserve privacy and verifiability. Our results demonstrate that this decentralized, quantum-secure system enables the deployment of privacy-preserving AI in banking in a manner that is robust, compliant with regulations, and prepared for the post-quantum era.

Index Terms—component, formatting, style, styling, insert

I. INTRODUCTION

Federated Learning (FL) has emerged as a leading paradigm for training machine learning models on decentralized data, particularly in privacy-sensitive domains such as banking and healthcare. By training models locally on edge devices and only communicating model updates, FL avoids the need to centralize raw user data, mitigating significant privacy risks. The most common approach, Federated Averaging (FedAvg), involves a central server that orchestrates the learning process by distributing a global model to clients, who then return their locally trained updates for aggregation.

However, the communication channel in FL presents a critical security vulnerability. The model parameters, exchanged between clients and the server, are susceptible to eavesdropping and malicious tampering from outside adversaries. While classical cryptographic methods like RSA and ECC have been used to secure this channel, they face an existential threat from

the advancement of quantum computing. Shor’s algorithm, for instance, can efficiently solve the mathematical problems underlying these cryptosystems, rendering them insecure.

To counter this looming threat, the field of Post-Quantum Cryptography (PQC) has developed new algorithms designed to be secure against attacks from both classical and quantum computers. The U.S. National Institute of Standards and Technology (NIST) has led a multi-year effort to standardize PQC algorithms, resulting in the selection of schemes like CRYSTALS-Kyber for key exchange and CRYSTALS-Dilithium and Falcon for digital signatures.

This paper introduces a novel framework that integrates PQC into a decentralized FL system to secure credit risk analysis. Our primary contribution is a comprehensive PQC-FL architecture that ensures end-to-end security without compromising model utility. We implement a hybrid cryptographic protocol using Kyber, Dilithium, Falcon, and AES to provide confidentiality, integrity, and authenticity. Furthermore, we employ a decentralized three-server architecture to enhance system robustness and fault tolerance. We empirically evaluate our framework using a non-IID credit risk dataset, demonstrating that it effectively thwarts malicious tampering attacks and provides future-proof security with quantifiable and manageable performance trade-offs.

II. THREATS AND OBJECTIVES

A. Threat Model

Our work aims to defend the FL process against an outside adversary capable of conducting the following attacks:

- **Man-in-the-Middle (MITM) / Poisoning Attack:** The primary threat we address is an outsider poisoning attack, where an adversary intercepts the communication channel between a client and the server. The attacker aims to inject malicious or corrupted model updates to degrade the global model’s performance or introduce targeted backdoors. This compromises the integrity of the training process.
- **Eavesdropping Attack:** An adversary could passively monitor the communication channel to capture model updates. Although these updates are not raw data, they can still leak sensitive information through model inversion or membership inference attacks. This attack compromises the confidentiality of the model parameters.

- **The Quantum Threat:** This is the overarching threat that motivates our work. A sophisticated adversary with access to a large-scale quantum computer could break the classical public-key cryptography (e.g., RSA, ECC) typically used to protect communication channels. This would render classical digital signatures forgeable and key exchanges insecure, nullifying defenses against the aforementioned attacks.

B. Objectives

To counter these threats, we defined the following research objectives:

- **Integrate and Benchmark PQC in FL:** To develop a robust FL framework enhanced with PQC using the Flower library and liboqs, and to establish baseline performance metrics for a realistic credit risk use case on non-IID data.
- **Design and Implement Hybrid Security:** To design and implement a hybrid security protocol that combines a PQC Key Encapsulation Mechanism (KEM) for confidentiality with an adaptive PQC digital signature scheme for integrity and authenticity, optimizing security and performance.
- **Quantify and Analyze Trade-offs:** To systematically measure and analyze the impact of integrating PQC on key performance indicators, including model accuracy, training time, communication cost, and the quality of aggregation, particularly in a non-IID environment.

III. METHODOLOGY

To achieve our objectives, we designed a decentralized, PQC-enhanced Federated Learning framework. The architecture and workflow are detailed below.

A. Decentralized Federated Learning architecture

Our framework diverges from the traditional centralized FL model. Instead, it employs a decentralized architecture consisting of multiple Sub-Aggregate Servers and one Global Aggregation Server, as shown in the architecture diagram.

- **Client Assignment:** Each client is assigned to one of the sub-aggregate servers.
- **Local Aggregation:** The sub-servers are responsible for receiving encrypted model updates from their assigned clients, performing PQC decryption and signature verification, and then securely aggregating the valid models.
- **Global Aggregation:** The Global Aggregation Server receives the locally aggregated models from each sub-server and performs a final aggregation to produce the updated global model for the next round. This hierarchical design enhances fault tolerance and reduces the risk of a single point of failure.

B. Hybrid PQC Security Layer

The core of our framework is a PQC security layer that protects the model updates during transit. This layer is implemented as a wrapper that intercepts communications between clients and servers.

- 1) **Confidentiality via Kyber KEM and AES:** We use a hybrid encryption scheme to ensure the confidentiality of model updates.
- 2) The server generates a Kyber512 public-private key pair and distributes the public key.
- 3) A client uses this public key to perform a Kyber512 key encapsulation, generating a shared secret and a ciphertext.
- 4) This shared secret is then used to derive a symmetric key for the AES algorithm (implemented via the Fernet library), which efficiently encrypts the large model parameter update.
- 5) The client sends both the Kyber ciphertext and the AES-encrypted model update to its assigned sub-server, which uses its Kyber512 private key to decrypt the payload.
- 6) **Integrity and Authenticity via Adaptive Digital Signatures:** To ensure integrity and authenticity, we use NIST-standardized PQC digital signatures.
- 7) We employ an adaptive strategy using Dilithium2 and Falcon-512. Clients with larger data contributions use Falcon-512 (which offers smaller signatures), while others use Dilithium2, striking a balance between performance and communication overhead.
- 8) After encrypting its update, each client signs the encrypted payload using its private signature key.
- 9) Upon receipt, the sub-server uses the client's public signature key to verify the signature before proceeding with decryption. If verification fails, the update is rejected as it is considered tampered or inauthentic.

C. Malicious Attack Simulation

To empirically validate the effectiveness of our signature scheme, we simulate an MITM attack. One designated client in the simulation is programmed to be malicious. This client follows the protocol correctly to generate an encrypted and signed update. However, after signing, it deliberately tampers with the bytes of the encrypted payload before transmission. This directly tests the server's ability to detect a signature mismatch and reject the corrupted update.

IV. EXPERIMENTS

Our evaluation was conducted in a simulated environment designed to rigorously test our framework's performance and security under realistic conditions.

A. Experimental Setup

The federated learning simulation was orchestrated using the Flower framework. The core Post-Quantum Cryptography functionalities were implemented using the liboqs-python library, which provides efficient bindings to the C implementations of the NIST-standardized PQC algorithms. The neural network was built and trained in PyTorch, with data handling managed by Pandas and Scikit-learn.

B. Dataset and Model

We utilized a public Indian credit risk dataset, partitioned among 50 clients in a non-IID manner to create a natural skew in data distribution and quantity. To handle the mixed data types, we designed a CreditRiskNet model featuring nn.Embedding layers for categorical features and nn.BatchNorm1d for numerical features, followed by dense layers with ReLU activations and a final Sigmoid output. The key hyperparameters used for training are detailed in Table I.

TABLE I
FEDERATED LEARNING HYPERPARAMETERS

Parameter	Value	Parameter	Value
Framework	Flower	Local Epochs per Client	3
Number of Clients	50	Batch Size	32
Number of Rounds	15-20	Optimizer	AdamW
Loss Function	BCELoss	Learning Rate	0.01

C. PQC Algorithm and Characteristics

The selection of PQC algorithms was central to our design. Our hybrid security layer uses Kyber512 for key encapsulation and an adaptive strategy for digital signatures based on the characteristics outlined in Table II. We chose Dilithium2 and Falcon-512 as they offer a compelling balance of security, performance, and size. The table compares them against the classical ECDSA-P256 algorithm to highlight the inherent overheads of PQC.

TABLE II
COMPARISON OF DIGITAL SIGNATURE ALGORITHM CHARACTERISTICS

Metric	Dilithium2	Falcon-512	ECDSA-P256
Security Basis	Lattice-based	Lattice-based	Elliptic Curve
Public Key Size (bytes)	1312	897	~65
Private Key Size (bytes)	2528	1281	~32
Signature Size (bytes)	2420	666	~72
Signing Time (ms)	~2.9	~5.0	~0.05
Verification Time (ms)	~0.9	~0.45	~0.1

V. RESULTS AND DISCUSSION

Our experiments provide a multi-faceted evaluation of the PQC-FL framework, culminating in key findings regarding security, model performance, and the inherent overheads of post-quantum cryptography. The results, summarized in Table III, paint a clear picture: our framework successfully establishes a quantum-resistant security posture while preserving the core machine learning utility of the federated system. While model accuracy remains stable across all tested scenarios, this security is achieved at the cost of quantifiable increases in computational and communication overhead, confirming the central trade-offs of implementing PQC.

A. Model Utility and Learning Process

A crucial outcome of our study is the confirmation that integrating PQC does not degrade the model's ability to learn. Figure 1 provides compelling visual evidence of this,

TABLE III
OVERALL PERFORMANCE AND SECURITY SUMMARY

Metric	PQC-FL	Classical (Sim.)	No Cryptography
Final Model Accuracy	88%	88%	88%
Malicious Update Detected	Yes	No	No
Avg. Compute Overhead (s)	10^{-2}	10^{-3}	N/A
Avg. Comm. Cost / Client (MB)	0.14	0.10	0.10




Fig. 1. Model accuracy comparison across 15 rounds, showing that PQC integration does not degrade learning performance.

where the accuracy curves of the PQC-FL, classical, and non-cryptographic scenarios are tightly interwoven throughout the 15 training rounds. This near-identical performance trajectory demonstrates that our cryptographic layer functions as a true security wrapper—it secures the data in transit without altering the mathematical integrity of the model updates. Consequently, the Federated Averaging algorithm's convergence is not hindered, allowing the global model to successfully achieve a high accuracy of approximately 88% in all scenarios. This result is a critical validation that future-proof security can be achieved without sacrificing the primary goal of the collaborative learning task.

B. Security Effectiveness and integrity verification

The primary security validation of our framework is demonstrated in Figure 2. This plot provides definitive proof of our system's defense against data tampering. In the round where the malicious client attempted to inject a corrupted update, the PQC-enabled server's verification mechanism correctly identified the mismatch between the digital signature and the altered payload. This cryptographic failure immediately triggered the rejection of the update, preventing the poisoned parameters from compromising the global model. In contrast, the baseline scenarios lacking this PQC signature verification would have




Fig. 2. The framework’s integrity check. The single bar shows a malicious update being detected and rejected by the PQC-FL system, while non-PQC methods would fail to do so.

unknowingly accepted the malicious data, highlighting their vulnerability to such outsider attacks. This result empirically validates the integrity and authenticity guarantees provided by our PQC layer.

C. Performance Overheads and Trade-offs

Achieving quantum resistance comes with inherent performance costs, which our experiments clearly quantify. **Computational Overhead:** Figure 3 illustrates the significant computational demands of PQC. The use of a logarithmic y-axis is necessary to effectively visualize the performance gap: the PQC operations consistently introduce an overhead that is an order of magnitude greater than that of simulated classical cryptography. This increased latency stems from the computationally intensive nature of the underlying lattice-based mathematics, such as large polynomial arithmetic, which is fundamentally more complex than the operations used in classical ECC. **Communication Overhead:** The higher communication cost summarized in Table III is a direct consequence of the larger key and signature sizes detailed in Table II. A Dilithium2 signature, at 2420 bytes, is over 30 times larger than a classical ECDSA signature (72 bytes). This substantial increase means that more data must be transmitted over the network in each round, leading to higher bandwidth consumption. Our adaptive signature strategy, which uses the more compact Falcon-512 signatures for clients with larger data contributions, is a pragmatic design choice that directly addresses this trade-off by minimizing the communication burden from the most impactful clients.

VI. CONCLUSION

Our research provides a definitive validation of the proposed Decentralized PQC-FL framework, proving its capability to deliver quantum-resistant security for sensitive AI workloads like credit risk analysis. The empirical results lead to several assertive conclusions:

- **Uncompromised Model Utility:** The integration of PQC security was achieved with 0% degradation in final model




Fig. 3. Computational overhead per round on a logarithmic scale, highlighting the performance cost of PQC.

accuracy. Our framework consistently converged to approximately 88% accuracy, a performance statistically identical to scenarios with classical or no cryptography, demonstrating that the security layer is non-invasive to the core learning process.

- **Absolute Security Against Tampering:** The framework demonstrated 100% effectiveness in detecting and rejecting the simulated outsider attack. The PQC digital signature verification provided a robust integrity check that is entirely absent in non-cryptographic systems, confirming its necessity for building a truly secure FL ecosystem.
- **Quantified Performance Trade-offs:** This enhanced security comes at a precisely quantified cost. Computationally, PQC operations introduced an overhead an order of magnitude (approximately 1000%) greater than classical methods in our simulation. The impact on communication is even more pronounced; the *Dilithium2* signature (2,420 bytes) is over 3300% larger than its classical ECDSA counterpart (72 bytes). This directly contributed to a measured 40% increase in data transfer per client in our specific FL setup.

In essence, our work transitions the concept of PQC-FL from a theoretical defense to a practical reality with clearly defined performance characteristics. We have proven that a future-proof, privacy-preserving AI ecosystem is not only possible but can be achieved with manageable and well-understood trade-offs, providing a validated blueprint for next-generation secure AI.

VII. FUTURE WORK

Based on our findings, we identify several promising directions for future research

- **Optimize PQC Performance:** Explore hardware acceleration for PQC operations on edge devices and dynamic algorithm selection based on real-time network conditions and compute capabilities.
- **Strengthen Privacy Guarantees:** Integrate additional privacy-enhancing technologies like Differential Privacy

or Homomorphic Encryption to provide defense-in-depth against a wider range of inference attacks.

- Improve Robustness: Incorporate Byzantine fault-tolerant aggregation rules and advanced malicious client detection mechanisms to defend against sophisticated insider attacks.
- Enhance Scalability: Implement dynamic server management and explore peer-to-peer FL architectures to improve scalability for massive-scale deployments.
- Expand Applicability: Adapt the framework to support more complex models, such as Large Language Models (LLMs), and apply it to other sensitive domains like healthcare for medical imaging analysis.

VIII. REFERENCES

- 1) P. Li, T. Chen, and J. Liu, "Enhancing Quantum Security over Federated Learning via Post-Quantum Cryptography,"
- 2) arXiv preprint arXiv:2409.04637, 2024.
- 3) H. B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, "Communication-efficient learning of deep networks from decentralized data," in
- 4) Proceedings of the 20th International Conference on Artificial Intelligence and Statistics (AISTATS), 2017.
- 5) L. Ducas, E. Kiltz, T. Lepoint, V. Lyubashevsky, P. Schwabe, G. Seiler, and D. Stehle, "CRYSTALS-Kyber: a CCA-secure module-lattice-based KEM," in 2018 IEEE European Symposium on Security and Privacy (EuroS&P), 2018, pp. 312-327.
- 6) L. Ducas, T. Lepoint, V. Lyubashevsky, P. Schwabe, G. Seiler, and D. Stehle, "CRYSTALS-Dilithium: A Lattice-Based Digital Signature Scheme," IACR Transactions on Cryptographic Hardware and Embedded Systems, vol. 2018, no. 1, pp. 238–268, Feb. 2018.
- 7) T. Prest, P. Fouque, J. Hoffstein, P. Kirchner, J. P. Silverman, and W. Whyte, "Falcon: Fast-fourier lattice-based compact signatures over NTRU," NIST Post-Quantum Cryptography Standardization Project, Round 1 Submission, 2017.
- 8) P. W. Shor, "Polynomial-time algorithms for prime factorization and discrete logarithms on a quantum computer," SIAM Journal on Computing, vol. 26, no. 5, pp. 1484-1509, Oct. 1997.
- 9) R. L. Rivest, A. Shamir, and L. Adleman, "A method for obtaining digital signatures and public-key cryptosystems," Communications of the ACM, vol. 21, no. 2, p. 120-126, feb 1978.
- 10) W. Wang et al., "FedCredit: A Secure and Privacy-Preserving Federated Learning Framework for Credit Scoring," [Journal/Conference, Year]. (Placeholder based on literature survey)
- 11) B. Commey, F. J. P. F. de Oliveira, R. A. L. Rabelo, and P. B. S. Lizcano, "PQS-BFL: A Post-Quantum Secure Blockchain-based Federated Learning," [Journal/Conference, Year]. (Placeholder based on literature survey)